

RUOHAN GAO

University of Maryland, College Park
Department of Computer Science
8125 Paint Branch Dr
College Park, MD 20742

Office: IRB-4248
E-mail: rhgao@umd.edu
<https://ruohangao.github.io/>

CURRENT APPOINTMENT

University of Maryland, College Park, College Park, MD
Assistant Professor at Department of Computer Science

Jan. 2025 – Present

EDUCATION

The University of Texas at Austin, Austin, TX
Ph.D. in Computer Science, Jan. 2021

The Chinese University of Hong Kong, Hong Kong
B.Eng. in Information Engineering, *First Class Honours*, July 2015

APPOINTMENTS

University of Maryland, College Park, College Park, MD
Adjunct Assistant Professor at Department of Computer Science

Feb. 2024 – Dec. 2024

Meta Reality Labs Research, Redmond, WA
Research Scientist

June 2023 – Dec. 2024

Stanford University, Stanford, CA
SAIL Postdoctoral Research Fellow

Feb. 2021 – June 2023

Facebook AI Research (FAIR), Austin, TX
Visiting Researcher

March 2020 – Dec. 2020

The University of Texas at Austin, Austin, TX
Graduate Student Researcher

Aug. 2015 – Jan. 2021

Facebook AI Research (FAIR), Cambridge, MA
Research Intern

June 2019 – Sept. 2019

Facebook AI Research (FAIR), Menlo Park, CA
Research Intern

May 2018 – Aug. 2018

HONORS AND AWARDS

- Best Paper Award, Tactile Sensing for Robotic Foundation Models Workshop, RSS 2026
For the paper “OmniTacTune: Policy-Agnostic Real-World RL for Tactile Residual Adaptation of Visual Policies” with Kelin Yu, Haode Zhang, Harish Ravichandar and Yunhai Han
- Sony Faculty Innovation Award, 2025
- Best Student Paper Award, IEEE WASPAA, 2025
For the paper “Towards Perception-Informed Latent HRTF Representations” with You Zhang, Andrew Francl, Paul Calamia, Zhiyao Duan, and Ishwarya Ananthabhotla
- AAAI New Faculty Highlights, 2025
- Stanford Artificial Intelligence Lab (SAIL) Postdoctoral Fellowship, 2021-2023

- Best Paper Award Runner Up, British Machine Vision Conference (BMVC), 2021
For the paper “Geometry-Aware Multi-Task Learning for Binaural Audio Generation from Video” with R. Garg and K. Grauman
- University Nominee for ACM Doctoral Dissertation Award, UT Austin, 2021.
- Michael H. Granof Award, UT Austin, 2021
UT Austin’s top graduate student award, awarded to top 1 dissertation in all areas from all graduate students of 2021
- Outstanding Dissertation Award in Mathematics, Engineering, Physical Sciences, and Biological and Life Sciences, UT Austin, 2021
- Google PhD Fellowship, 2019 - 2021
- Outstanding Reviewer Award, CVPR 2020, NeurIPS 2024
- Graduate Dean’s Prestigious Fellowship Supplement Fellow, UT Austin, 2019 & 2020
- Adobe Research Fellowship, 2019
- Best Paper Award Finalist, Conference on Computer Vision and Pattern Recognition (CVPR) 2019
For the paper, “2.5D Visual Sound”, with K. Grauman
- Sir Run Run Shaw Postgraduate Scholarship, CUHK, 2015
- Dean’s List, Engineering Faculty, CUHK, 2014 & 2015
- Cheung Wang Ngai Joseph GOAL Programme Scholarship, CUHK, 2014
- Q W Lee Scholarship, Top Academic Excellence Scholarship, CUHK, 2014

PEER-REVIEWED CONFERENCE PAPERS

(* indicates equal contribution; † indicates equal advising.)

1. Kelin Yu*, Haode Zhang*, Harish Ravichandar, Yunhai Han, **Ruohan Gao**, “OmniTacTune: Policy-Agnostic Real-World RL for Tactile Residual Adaptation of Visual Policies”, in *Proceedings of the Conference on Robot Learning (CoRL)*, Nov. 2026.
2. Dongyu Luo*, Kelin Yu*, Amir-Hossein Shahidzadeh, Cornelia Fermuler, Yiannis Aloimonos, **Ruohan Gao**, “ControlTac: Scaling Tactile Data with Physically Controlled Tactile Image Generation”, in *Proceedings of the Conference on Robot Learning (CoRL)*, Nov. 2026.
3. Changyeon Lee, Kelin Yu, Zheng Wei, Shuna Ni, **Ruohan Gao**, “OdorSpaces: Visual-Olfactory Embodied Navigation in 3D Environments”, in *Proceedings of the Conference on Robot Learning (CoRL)*, Nov. 2026.
4. Zisen Shao*, Zihao Wei*, Derong Jin and **Ruohan Gao**, “Objects as Audio-Visual Model Sound Fields”, in *Proceedings of the European Conference on Computer Vision (ECCV)*, Sept. 2026.
5. Derong Jin*, Xiyi Chen*, Ming Lin, and **Ruohan Gao**, “SonoWorld: From One Image to a 3D Audio-Visual Scene”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2026.
6. Ramaneswaran S, Koushik Jayakumar, S Sakshi, Sreyan Ghosh, **Ruohan Gao**, Dinesh Manocha, “Do Audio-Visual Large Language Models Really See and Hear?”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR Findings)*, June 2026.
7. Derong Jin and **Ruohan Gao**, “Differentiable Room Acoustic Rendering with Multi-View Vision Priors”, in *Proceedings of the International Conference on Computer Vision (ICCV)*, Oct. 2025. (**Oral presentation**, 0.5% acceptance rate)
8. Sanjoy Chowdhury, Subrata Biswas, Sayan Nag, Tushar Nagarajan, Calvin Murdock, Ishwarya Ananthabhotla, Yijun Qian, Vamsi Krishna Ithapu, Dinesh Manocha, and **Ruohan Gao**, “EgoAdapt: Adaptive Multisensory Distillation and Policy Learning for Efficient Egocentric Perception”, in *Proceedings of the International Conference on Computer Vision (ICCV)*, Oct. 2025.

9. Kelin Yu, Sheng Zhang, Harshit Soora, Furong Huang, Heng Huang, Pratap Tokekar, and **Ruohan Gao**, “GenFlowRL: Shaping Rewards with Generative Object-Centric Flow in Visual Reinforcement Learning”, in *Proceedings of the International Conference on Computer Vision (ICCV)*, Oct. 2025.
10. Sanjoy Chowdhury*, Sayan Nag*, Subhrajyoti Dasgupta*, Yaoting Wang, Mohamed Elhoseiny, **Ruohan Gao**, and Dinesh Manocha, “AVTrustBench: Assessing and Enhancing Reliability and Robustness in Audio-Visual LLMs”, in *Proceedings of the International Conference on Computer Vision (ICCV)*, Oct. 2025.
11. Sanjoy Chowdhury*, Hanan Gani*, Nishit Anand, Sayan Nag, **Ruohan Gao**, Mohamed Elhoseiny, Salman Khan, and Dinesh Manocha, “Aurelia: Test-time Reasoning Distillation in Audio-Visual LLMs”, in *Proceedings of the International Conference on Computer Vision (ICCV)*, Oct. 2025.
12. You Zhang, Andrew Franci, **Ruohan Gao**, Paul Calamia, Zhiyao Duan, and Ishwarya Ananthabhotla, “Towards Perception-Informed Latent HRTF Representations” in *Proceedings of IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, Oct. 2025. **[Best Student Paper Award]**
13. Ricardo Falcon Perez, **Ruohan Gao**, Gregor Mueckl, Sebastia V. Amengual Gari, Ishwarya Ananthabhotla, “Scene-wide Acoustic Parameter Estimation” in *Proceedings of IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, Oct. 2025.
14. Xiulong Liu, Anurag Kumar, Paul Calamia, Sebastia Vicenc Amengual Gari, Calvin Murdock, Ishwarya Ananthabhotla, Philip W Robinson, Eli Shlizerman, Vamsi Krishna Ithapu, **Ruohan Gao**, “Hearing Anywhere in Any Environment”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2025.
15. Chao Huang, **Ruohan Gao**, J. M. F. Tsang, Jan Kurcius, Cagdas Bilen, Chenliang Xu, Anurag Kumar, Sanjeel Parekh, “Learning to Highlight Audio by Watching Movies”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2025.
16. Heeseung Yun, **Ruohan Gao**, Ishwarya Ananthabhotla, Anurag Kumar, Jacob Donley, Chao Li, Gunhee Kim, Vamsi Krishna Ithapu, and Calvin Murdock, “Spherical World-Locking for Audio-Visual Localization in Egocentric Videos”, in *Proceedings of the European Conference on Computer Vision (ECCV)*, Sept. 2024.
17. Sanjoy Chowdhury*, Sayan Nag*, Subhrajyoti Dasgupta*, Jun Chen, Mohamed Elhoseiny†, **Ruohan Gao**†, and Dinesh Manocha†, “Meerkat: Audio-Visual Large Language Model for Grounding in Space and Time”, in *Proceedings of the European Conference on Computer Vision (ECCV)*, Sept. 2024.
18. Xutong Jin*, Chenxi Xu*, **Ruohan Gao**, Jiajun Wu, Guoping Wang, Sheng Li, “DiffSound: Differentiable Modal Sound Rendering and Inverse Rendering for Diverse Inference Tasks”, in *Proceedings of Special Interest Group on Computer Graphics and Interactive Techniques Conference (SIGGRAPH)*, July 2024
19. Mason L. Wang*, Ryosuke Sawata*, Samuel Clarke, **Ruohan Gao**, Shangzhe Wu, and Jiajun Wu, “Hearing Anything Anywhere”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2024.
20. Wenqi Jia, Miao Liu, Hao Jiang, Ishwarya Ananthabhotla, James M Rehg, Vamsi Krishna Ithapu, and **Ruohan Gao**, “The Audio-Visual Conversational Graph: From an Egocentric-Exocentric Perspective”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2024.
21. Mason Wang*, Samuel Clarke*, Jui-Hsien Wang, **Ruohan Gao**, and Jiajun Wu, “SoundCam: A Dataset for Tasks in Tracking and Identifying Humans from Real Room Acoustics”, in *Proceedings of the Conference on Neural Information Processing Systems Datasets and Benchmarks Track (NeurIPS)*, Dec. 2023.

22. Sharon Lee*, Ruohan Zhang*, Minjune Hwang*, Ayano Hiranaka*, Chen Wang, Wensi Ai, Jin Jie Ryan Tan, Shreya Gupta, Yilun Hao, Gabrael Levine, **Ruohan Gao**, Anthony Norcia, Jiajun Wu†, and Li Fei-Fei†, “NOIR: Neural Signal Operated Intelligent Robot for Everyday Activities”, in *Proceedings of the Conference on Robot Learning (CoRL)*, Nov. 2023.
23. **Ruohan Gao***, Yiming Dou*, Hao Li*, Tanmay Agarwal, Jeannette Bohg, Yunzhu Li, Li Fei-Fei, Jiajun Wu, “The ObjectFolder Benchmark: Multisensory Object-Centric Learning with Neural and Real Objects”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2023.
24. Samuel Clarke, **Ruohan Gao**, Mason Wang, Mark Rau, Julia Xu, Jui-Hsien Wang, Doug James, Jiajun Wu, “RealImpact: A Dataset of Impact Sound Fields for Real Objects”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2023. (**Highlight paper**)
25. **Ruohan Gao***, Hao Li*, Gokul Dharan, Zhuzhu Wang, Chengshu Li, Fei Xia, Silvio Savarese, Li Fei-Fei, Jiajun Wu, “Sonicverse: A Multisensory Simulation Platform for Embodied Household Agents that See and Hear”, in *Proceedings of the International Conference on Robotics and Automation (ICRA)*, May 2023.
26. Trevor Scott Standley, **Ruohan Gao**, Dawn Chen, Jiajun Wu, Silvio Savarese, “An Extensible Multi-modal Multi-task Object Dataset with Materials”, in *Proceedings of the International Conference on Learning Representations (ICLR)*, May 2023.
27. Hao Li*, Yizhi Zhang*, Junzhe Zhu, Shaoxiong Wang, Michelle A Lee, Huazhe Xu, Edward Adelson, Li Fei-Fei, **Ruohan Gao†**, Jiajun Wu†, “See, Hear, and Feel: Smart Sensory Fusion for Robotic Manipulation”, in *Proceedings of the Conference on Robot Learning (CoRL)*, Dec. 2022.
28. **Ruohan Gao***, Zilin Si*, Yen-Yu Chang*, Samuel Clarke, Jeannette Bohg, Li Fei-Fei, Wenzhen Yuan, Jiajun Wu, “ObjectFolder 2.0: A Multisensory Object Dataset for Sim2Real Transfer”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2022.
29. Changan Chen, **Ruohan Gao**, Paul Calamia, Kristen Grauman, “Visual Acoustic Matching”, to appear in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2022. (**Oral presentation**, 4.2% acceptance rate)
30. **Ruohan Gao**, Yen-Yu Chang*, Shivani Mall*, Li Fei-Fei, Jiajun Wu, “ObjectFolder: A Dataset of Objects with Implicit Visual, Auditory, and Tactile Representations”, in *Proceedings of the Conference on Robot Learning (CoRL)*, Nov. 2021.
31. Samuel Clarke, Negin Heravi, Mark Rau, **Ruohan Gao**, Jiajun Wu, Doug James, Jeannette Bohg, “DiffImpact: Differentiable Rendering and Identification of Impact Sounds”, in *Proceedings of the Conference on Robot Learning (CoRL)*, Nov. 2021. (**Oral presentation**, 6.5% acceptance rate)
32. Rishabh Garg, **Ruohan Gao**, Kristen Grauman, “Geometry-Aware Multi-Task Learning for Binaural Audio Generation from Video”, in *Proceedings of the British Machine Vision Conference (BMVC)*, Nov. 2021. (**Oral presentation**, 3.3% acceptance rate) [**Best Paper Award Runner Up**]
33. **Ruohan Gao** and Kristen Grauman, “VisualVoice: Audio-Visual Speech Separation with Cross-Modal Consistency”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2021.
34. Changan Chen, Sagnik Majumder, Ziad Al-Halah, **Ruohan Gao**, Santhosh K. Ramakrishnan and Kristen Grauman, “Learning to Set Waypoints for Audio-Visual Navigation,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, May 2021.
35. **Ruohan Gao**, Changan Chen, Ziad Al-Halah, Carl Schissler, and Kristen Grauman, “VisualEchoes: Spatial Image Representation Learning through Echolocation”, in *Proceedings of the European Conference on Computer Vision (ECCV)*, Aug. 2020.

36. **Ruohan Gao**, Tae-Hyun Oh, Kristen Grauman, and Lorenzo Torresani, “Listen to Look: Action Recognition by Previewing Audio”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2020.
37. **Ruohan Gao** and Kristen Grauman, “Co-Separating Sounds of Visual Objects”, in *Proceedings of the International Conference on Computer Vision (ICCV)*, Oct. 2019.
38. **Ruohan Gao** and Kristen Grauman, “2.5D Visual Sound”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2019. (**Oral presentation**, 5.6% acceptance rate) [**Best Paper Award Finalist**]
39. **Ruohan Gao**, Rogerio Feris, and Kristen Grauman, “Learning to Separate Object Sounds by Watching Unlabeled Video”, in *Proceedings of the European Conference on Computer Vision (ECCV)*, Sept. 2018. (**Oral presentation**, 2.4% acceptance rate)
40. Dinesh Jayaraman, **Ruohan Gao**, and Kristen Grauman, “ShapeCodes: Self-Supervised Feature Learning by Lifting Views to Viewgrids”, in *Proceedings of the European Conference on Computer Vision (ECCV)*, Sept. 2018.
41. **Ruohan Gao**, Bo Xiong, and Kristen Grauman, “Im2Flow: Motion Hallucination from Static Images for Action Recognition”, in *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2018. (**Oral presentation**, 2.1% acceptance rate)
42. **Ruohan Gao** and Kristen Grauman, “On-demand Learning for Deep Image Restoration”, in *Proceedings of the International Conference on Computer Vision (ICCV)*, Oct. 2017.
43. **Ruohan Gao**, Dinesh Jayaraman, and Kristen Grauman, “Object-centric Representation Learning from Unlabeled Videos”, in *Proceedings of the Asian Conference on Computer Vision (ACCV)*, Nov. 2016.
44. **Ruohan Gao**, Huanle Xu, Pili Hu, and Wing Cheong Lau, “Accelerating Graph Mining Algorithms via Uniform Random Edge Sampling”, in *Proceedings of IEEE International Conference on Communications (ICC)*, May 2016.
45. **Ruohan Gao**, Pili Hu, and Wing Cheong Lau, “Property Preservation under Community-Based Sampling”, in *Proceedings of the Global Communications Conference (GLOBECOM)*, Dec. 2015.
46. **Ruohan Gao**, Huanle Xu, Pili Hu, and Wing Cheong Lau, “Accelerating Graph Mining Algorithms via Uniform Random Edge Sampling (Poster)”, in *ACM Conference on Online Social Networks (COSN)*, Nov. 2015.

PEER-REVIEWED JOURNAL ARTICLES

(* indicates equal contribution; † indicates equal advising.)

1. Neel Mokaria, Rishie Raj, Dheeraj Baiju, Xiaoqian Shen, Shraman Pramanick, Kevin Qinghong Lin, Arda Senocak, Mike Zheng Shou, Philip Torr, Mohamed Elhoseiny, Yapeng Tian, **Ruohan Gao**, Salman Khan, Sayan Nag, Sanjoy Chowdhury, Dinesh Manocha. “Foundations and Frontiers of Multimodal Agentic Frameworks”, in *Transactions on Machine Learning Research (TMLR)*, May 2026.
2. **Ruohan Gao**, “Multisensory Machine Intelligence”, in *AI Magazine*, Aug. 2025. (**Invited article for AAAI-25 New Faculty Highlights**)
3. Rishabh Garg, **Ruohan Gao**, Kristen Grauman, “Visually-Guided Audio Spatialization in Video with Geometry-Aware Multi-Task Learning”, in *International Journal of Computer Vision (IJCV)*, May 2023. (**Invited article for best papers of BMVC 2021**)
4. Hong-Xing Yu*, Michelle Guo*, Alireza Fathi, Yen-Yu Chang, Eric Ryan Chan, **Ruohan Gao**, Thomas Funkhouser, Jiajun Wu, “Learning Object-Centric Neural Scattering Functions for Free-Viewpoint Relighting and Scene Composition”, in *Transactions on Machine Learning Research (TMLR)*, April 2023.

5. Simon Le Cleac'h, Hong-Xing Yu, Michelle Guo, Taylor A. Howell, **Ruohan Gao**, Jiajun Wu, Zachary Manchester, Mac Schwager, “Differentiable Physics Simulation of Dynamics-Augmented Neural Objects”, in *Robotics and Automation Letters* (RA-L), March 2023

INVITED AND CONFERENCE TALKS

- Speech and Audio in the Northeast (SANE) at MIT
Invited Talk, Cambridge, MA Oct. 2026
- Wild3D Workshop
Invited Talk, ECCV Workshop, Malmö, Sweden Sept. 2026
- AVGenL Workshop
Invited Talk, ECCV Workshop, Malmö, Sweden Sept. 2026
- Gen4AVC Workshop
Invited Talk, ECCV Workshop, Malmö, Sweden Sept. 2026
- Johns Hopkins University
JSALT Summer Workshop Lecture, Baltimore, MD June 2026
- Sight and Sound
Invited Talk, CVPR Workshop, Denver, CO June 2026
- SMLM Workshop at Yale University
Invited Talk, New Haven, CT April 2026
- Laboratory for Telecommunication Sciences (LTS)
Invited Talk, College Park, MD March 2026
- Dolby Laboratories
Invited Talk, Smart Realities Summit, San Francisco, CA Feb. 2026
- University of Pennsylvania
Invited Talk, GRASP Seminar, Philadelphia, PA Jan. 2026
- RoboAcoustics Workshop
Invited Talk, ICRA Workshop, Atlanta, GA May 2025
- Maryland Robotic Center Research Symposium
Invited Talk, College Park, MD May 2025
- Stanford CS231n 10th Anniversary
Invited Guest Lecture, Stanford University, Stanford, CA May 2025
- AAAI New Faculty Highlights
New Faculty Highlight Talk, AAAI 2025, Philadelphia, PA Feb. 2025
- SIGGRAPH 2024
Oral Presentation, Denver, CO Aug. 2024
- Sight and Sound
Invited Talk, CVPR Workshop, Seattle, WA June 2024
- Tsinghua University
Invited Talk, Beijing, China May 2024
- Peking University
Invited Talk, Beijing, China May 2024
- University of Michigan
Invited Guest Lecture, Virtual April 2024
- Edge AI and Vision
Invited Talk, Virtual Oct. 2023
- University of Southern California
Invited Talk, Los Angeles, CA April 2023

- Cornell University
Invited Talk, Ithaca, NY March 2023
- University of Maryland, College Park
Invited Talk, College Park, MD March 2023
- Georgia Institute of Technology
Invited Talk, Atlanta, GA Feb. 2023
- Gaoling School of Artificial Intelligence
Invited Talk, Renmin University of China, Virtual May 2022
- Intelligent Sensing Winter School
Invited Talk, Queen Mary University of London, Virtual Dec. 2021
- Mitsubishi Electric Research Labs (MERL)
Invited Talk, MERL Seminar Series, Virtual Sept. 2021
- Special Session on Low-Power Computer Vision
Invited Talk, AICAS, Virtual June 2021
- Stanford University
Interactive Perception and Robot Learning Lab, Stanford, CA April 2021
- Stanford University
Stanford Vision and Learning Lab, Stanford, CA March 2021
- Facebook Reality Lab
Invited Talk, FRL Audio, Redmond, WA Dec. 2020
- The University of California, Berkeley
Invited Talk, Computer Vision Seminar, Berkeley, CA May 2020
- Massachusetts Institute of Technology
Invited Talk, Spoken Language Systems Group, Cambridge, MA April 2020
- The University of Texas at San Antonio
Invited Talk, AI Consortium Seminar Series, San Antonio, TX April 2020
- Massachusetts Institute of Technology
Invited Talk, Vision Seminar, Cambridge, MA Sept. 2019
- Adobe Research
Adobe Research Fellowship Ceremony, San Jose, CA Aug. 2019
- Google
Lightning Talk, PhD Fellowship Summit, Mountain View, CA July 2019
- Sight and Sound
Invited Talk, CVPR Workshop, Long Beach, CA June 2019
- Conference on Computer Vision and Pattern Recognition (CVPR), 2019
Oral Presentation, Long Beach, CA June 2019
- European Conference on Computer Vision (ECCV), 2018
Oral Presentation, Munich, Germany Oct. 2018
- Sight and Sound
Oral Presentation, CVPR Workshop, Salt Lake City, Utah June 2018
- Conference on Computer Vision and Pattern Recognition (CVPR), 2018
Oral Presentation, Salt Lake City, Utah June 2018

THESES

- Ruohan Gao. "Look and Listen: From Semantic to Spatial Audio-Visual Perception". Ph.D. Thesis, The University of Texas at Austin, 2021.

PATENTS

- US20210174817A1: Systems and Methods for Visually Guided Audio Separation, 2021.

TEACHING EXPERIENCE**University of Maryland, College Park**, College Park, MD

CMSC848M: Multimodal Computer Vision	Spring 2026
CMSC498K: Introduction to Multimodal Deep Learning	Fall 2025
CMSC848M: Multimodal Computer Vision	Spring 2025

Stanford University, Stanford, CA

CS231N: Deep Learning for Computer Vision (with Fei-Fei Li and Yunzhu Li)	Spring 2023
CS231N: Deep Learning for Computer Vision (with Fei-Fei Li and Jiajun Wu)	Spring 2022

PROFESSIONAL SERVICES**Area Chair**

International Conference on Learning Representations (ICLR), 2027
 Conference on Robot Learning (CoRL), 2026
 Conference on Neural Information Processing Systems (NeurIPS), 2026
 European Conference on Computer Vision (ECCV), 2026
 Conference on Computer Vision and Pattern Recognition (CVPR), 2026
 International Conference on Learning Representations (ICLR), 2026
 International Conference on 3D Vision (3DV), 2025
 IEEE International Conference on Computer Vision (ICCV), 2023

Senior Program Committee

AAAI Conference on Artificial Intelligence (AAAI), 2023–2025

Organizer

Co-Organizer, Workshop on Sight and Sound, CVPR 2026
 Co-Organizer, Workshop on Sight and Sound, CVPR 2025
 Co-Organizer, Workshop on Sight and Sound, CVPR 2024
 Co-Organizer, Workshop on AV4D: Visual Learning of Sounds in Spaces, ICCV 2023
 Co-Organizer, Workshop on Sight and Sound, CVPR 2023
 Co-Organizer, Workshop on Creative AI Across Modalities, AAAI 2023
 Co-Organizer, Workshop on AV4D: Visual Learning of Sounds in Spaces, ECCV 2022
 Co-Organizer, Workshop on Sight and Sound, CVPR 2022
 Co-Organizer, Workshop on Sight and Sound, CVPR 2021
 Lead Organizer, Workshop on Embodied Multimodal Learning, ICLR 2021
 Co-Organizer, Workshop on Sight and Sound, CVPR 2020

Conference Program Committee Member / Reviewer

IEEE Conference on Computer Vision and Pattern Recognition (CVPR)
 IEEE International Conference on Computer Vision (ICCV)
 European Conference on Computer Vision (ECCV)
 ACM International Conference on Computer Graphics and Interactive Techniques (SIGGRAPH)
 ACM SIGGRAPH Conference and Exhibition on Computer Graphics and Interactive Techniques in Asia (SIGGRAPH Asia)
 International Conference on Learning Representations (ICLR)
 Conference on Neural Information Processing Systems (NeurIPS)
 International Conference on Machine Learning (ICML)
 Conference on Robot Learning (CoRL)
 Robotics: Science and Systems (RSS)

IEEE International Conference on Robotics and Automation (ICRA)
 AAAI Conference on Artificial Intelligence (AAAI)
 British Machine Vision Conference (BMVC)
 Asian Conference on Computer Vision (ACCV)
 IEEE Winter Conference on Computer Vision (WACV)
 International Symposium on Mixed and Augmented Reality (ISMAR)

Workshop Program Committee Member / Reviewer

Workshop on T4V: Transformers for Vision, CVPR 2022
 Workshop on Large Scale Holistic Video Understanding, CVPR 2021
 Workshop on Learning from Unlabeled Videos, CVPR 2021
 Workshop on Self-Supervised Learning: Theory and Practice, NeurIPS 2020
 Workshop on Multi-Modal Video Analysis, ECCV 2020

Journal Reviewer:

IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)
 IEEE Transactions on Circuits and Systems for Video Technology (TCSVT)
 IEEE Transactions on Image Processing (TIP)
 IEEE Transactions on Multimedia (TMM)
 IEEE/ACM Transactions on Audio Speech and Language Processing (TASLP)
 IEEE Transactions on Robotics (T-RO)
 IEEE Journal of Selected Topics in Signal Processing (J-STSP)
 Journal of Artificial Intelligence Research (JAIR)
 Computer Vision and Image Understanding (CVIU)
 Journal of the Acoustical Society of America (JASA)
 ACM Computing Surveys

SELECTED MEDIA COVERAGE

Stanford HAI , Wearable device allows humans to control robots with brain waves.	Jan. 2024
TechXplore , A multisensory simulation platform to train and test home robots.	June 2023
ENGADGET , Auditory AIs promise a more immersive AR/VR experience.	June 2022
Meta AI , Introducing AI-driven acoustic synthesis for AR and VR.	June 2022
The University of Texas at Austin , Looking and Listening in Machine Learning.	Nov. 2021
Facebook AI Blog , New milestones in embodied AI.	Aug. 2020
MIT Technology Review , Deep learning turns mono recordings into immersive sound.	Dec. 2018
Two Minute Papers , This AI produces binaural (2.5D) audio.	Jan. 2019
Facebook AI Blog , Creating 2.5D visual sound for an immersive audio experience.	June 2019